

# 05 - Aggregation

ST 597 | Spring 2017

University of Alabama

05-aggregate.pdf

## Contents

0.1	Required Packages and Data . . . . .	1
<b>1</b>	<b>Single table verbs</b>	<b>2</b>
1.1	dplyr single table verbs . . . . .	2
1.2	summarize() . . . . .	2
1.3	summary() is not summarize() . . . . .	3
<b>2</b>	<b>Group-wise operations</b>	<b>4</b>
2.1	Split-Apply-Combine . . . . .	4
2.2	group_by() . . . . .	5
2.3	Counting . . . . .	6
2.4	Chaining . . . . .	10
2.5	Multiple grouping levels . . . . .	11
2.6	Grouped Mutate and Filter . . . . .	13
2.7	Window Functions . . . . .	14

## 0.1 Required Packages and Data

```
library(tidyverse)
library(nycflights13)
data(flights)
```

# 1 Single table verbs

## 1.1 dplyr single table verbs

- `filter()`: select rows
- `arrange()`: reorder rows
  - `desc()` to use descending order
- `select()`: select certain columns
  - helper functions: `starts_with()`, `ends_with()`, `matches()`, `contains()`, `?select`
- `mutate()`: modify or create new variables
  - `transmute()`: only return new variables
- `summarise()`: reduce variables to values
  - Most useful when data is grouped

## 1.2 `summarize()`

The `summarize()` function calculates summary statistics for a column (or multiple columns). It collapses a data frame to a *single row*:

```
summarize(flights, avg.dist = mean(distance)) # mean distance
#> # A tibble: 1 × 1
#>   avg.dist
#>   <dbl>
#> 1    1040

summarize(flights, avg.dist = mean(distance), med.dist = median(distance))
#> # A tibble: 1 × 2
#>   avg.dist med.dist
#>   <dbl>    <dbl>
#> 1    1040     872

summarize(flights,
  n.records = n(), # number of records
  n.missing = sum(is.na(arr_delay)), # number of NA's
  num.delay = sum(arr_delay>0, na.rm=TRUE), # num of delayed flights
  prop.delay = mean(arr_delay>0, na.rm=TRUE) ) # proportion of delayed flights
#> # A tibble: 1 × 4
#>   n.records n.missing num.delay prop.delay
#>   <int>    <int>    <int>    <dbl>
#> 1   336776     9430   133004     0.406
```


Just to check the numbers,  $133004 / (336776 - 9430) = 0.406$

It works like this `summarize(<data>, name1 = f(<colname>), name2=g(<colname>))` where *f, g* are some functions (e.g., `mean()`, `median()`, `sd()`).

It can also use functions that take more than one column as input, but most return a single value. E.g.,

```
#- correlation of dep_delay and arr_delay
summarize(flights, delay_cor = cor(dep_delay, arr_delay, use="complete.obs"))
#> # A tibble: 1 × 1
```

```
#> delay_cor
#> <dbl>
#> 1 0.915
# same as:
# cor(flights$dep_delay, flights$arr_delay, use="complete.obs")
ggplot(sample_n(flights, 10000), aes(dep_delay, arr_delay)) +
  geom_smooth() + geom_point(alpha=.3)
#> `geom_smooth()` using method = 'gam'
#> Warning: Removed 281 rows containing non-finite values (stat_smooth).
#> Warning: Removed 281 rows containing missing values (geom_point).
```



### 1.3 `summary()` is not `summarize()`

The `summary()` function is a base R function that reports some basic summary stats for all columns.

```
summary(flights) # base R function
```

	year	month	day	dep_time
#> Min.	:2013	Min. : 1.00	Min. : 1.0	Min. : 1
#> 1st Qu.:	:2013	1st Qu.: 4.00	1st Qu.: 8.0	1st Qu.: 907
#> Median :	:2013	Median : 7.00	Median :16.0	Median :1401
#> Mean :	:2013	Mean : 6.55	Mean :15.7	Mean :1349
#> 3rd Qu.:	:2013	3rd Qu.:10.00	3rd Qu.:23.0	3rd Qu.:1744
#> Max. :	:2013	Max. :12.00	Max. :31.0	Max. :2400
#>				NA's :8255

	sched_dep_time	dep_delay	arr_time	sched_arr_time
#> Min.	: 106	Min. : -43	Min. : 1	Min. : 1
#> 1st Qu.:	: 906	1st Qu.: -5	1st Qu.:1104	1st Qu.:1124
#> Median :	:1359	Median : -2	Median :1535	Median :1556
#> Mean :	:1344	Mean : 13	Mean :1502	Mean :1536
#> 3rd Qu.:	:1729	3rd Qu.: 11	3rd Qu.:1940	3rd Qu.:1945
#> Max. :	:2359	Max. :1301	Max. :2400	Max. :2359
#>		NA's :8255	NA's :8713	

```

#>   arr_delay   carrier      flight      tailnum
#> Min.      : -86   Length:336776   Min.      : 1   Length:336776
#> 1st Qu.: -17   Class :character   1st Qu.: 553   Class :character
#> Median : -5    Mode  :character   Median :1496   Mode  :character
#> Mean      : 7                                Mean      :1972
#> 3rd Qu.: 14                                3rd Qu.:3465
#> Max.      :1272                             Max.      :8500
#> NA's      :9430
#>   origin      dest      air_time      distance
#> Length:336776   Length:336776   Min.      : 20   Min.      : 17
#> Class :character Class :character   1st Qu.: 82   1st Qu.: 502
#> Mode  :character Mode  :character   Median :129   Median : 872
#> Mean      :151   Mean      :1040
#> 3rd Qu.:192   3rd Qu.:1389
#> Max.      :695   Max.      :4983
#> NA's      :9430
#>   hour      minute      time_hour
#> Min.      : 1.0   Min.      : 0.0   Min.      :2013-01-01 05:00:00
#> 1st Qu.: 9.0   1st Qu.: 8.0   1st Qu.:2013-04-04 13:00:00
#> Median :13.0   Median :29.0   Median :2013-07-03 10:00:00
#> Mean      :13.2   Mean      :26.2   Mean      :2013-07-03 05:02:36
#> 3rd Qu.:17.0   3rd Qu.:44.0   3rd Qu.:2013-10-01 07:00:00
#> Max.      :23.0   Max.      :59.0   Max.      :2013-12-31 23:00:00
#>

```

The `summarize()` function applies a function that summarizes each column down to a single number.

```

summarize(flights,
  min=min(arr_delay, na.rm=TRUE),
  Q1 = quantile(arr_delay, 0.25, na.rm=TRUE),
  median=median(arr_delay, na.rm=TRUE),
  mean = mean(arr_delay, na.rm=TRUE),
  Q3 = quantile(arr_delay, 0.75, na.rm=TRUE),
  max = max(arr_delay, na.rm=TRUE),
  count.NA = sum(is.na(arr_delay)))
#> # A tibble: 1 × 7
#>   min    Q1 median  mean    Q3    max count.NA
#>   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl>   <int>
#> 1   -86   -17     -5    6.9    14   1272     9430

```

## 2 Group-wise operations

### 2.1 Split-Apply-Combine

These operations are more powerful when they can be used with grouping variables. Split - Apply - Combine.

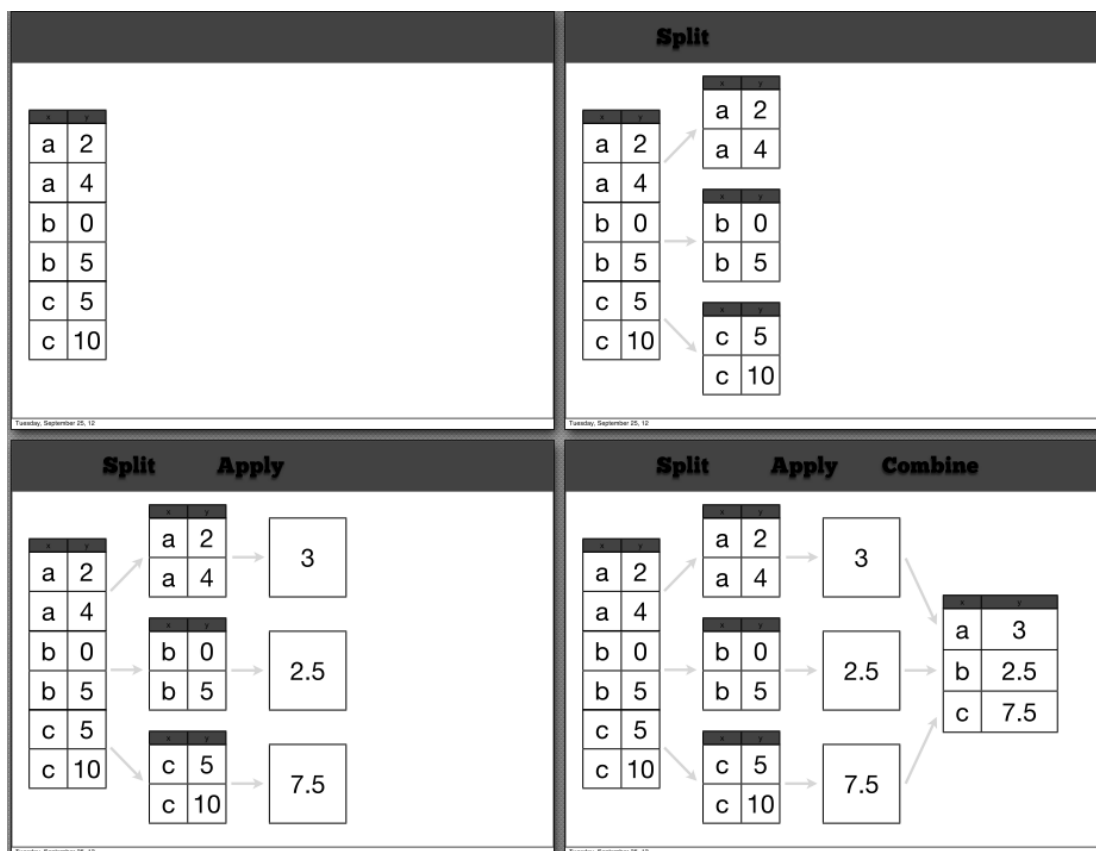


Image from Hadley Wickham UseR tutorial June 2014 <http://www.dropbox.com/sh/i8qnluwmuicxc/AAAgT9tIKoIm7WZKlyK25lh6a>

## 2.2 group\_by()

First use the `group_by()` function to group the data (determines how to split), then apply function(s) to each group using the `summarize()` function. Note: grouping should to be applied on discrete variables (categorical, factor, or maybe integer valued columns).

```
#- Get maximum delay by origin
by_origin = group_by(flights, origin)
summarize(by_origin, max.delay = max(arr_delay, na.rm=TRUE))
#> # A tibble: 3 × 2
#>   origin max.delay
#>   <chr>     <dbl>
#> 1 EWR      1109
#> 2 JFK      1272
#> 3 LGA       915
```

```
#- Get delay info by origin and destination
by_dest = group_by(flights, origin, dest)
summarize(by_dest,
  max.delay = max(arr_delay, na.rm=TRUE),
  avg.delay = mean(arr_delay, na.rm=TRUE),
```

```

    min.delay = min(arr_delay, na.rm=TRUE),
    count = n() )
#> Source: local data frame [224 x 6]
#> Groups: origin [?]
#>
#>   origin dest max.delay avg.delay min.delay count
#>   <chr> <chr>    <dbl>    <dbl>    <dbl> <int>
#> 1   EWR  ALB      328     14.397     -34    439
#> 2   EWR  ANC       39     -2.500     -47     8
#> 3   EWR  ATL     796     13.233     -39   5022
#> 4   EWR  AUS     349     -0.474     -59   968
#> 5   EWR  AVL     228      8.805     -26   265
#> 6   EWR  BDL     266      7.049     -43   443
#> 7   EWR  BNA     364     12.708     -41  2336
#> 8   EWR  BOS     422      4.784     -47  5327
#> 9   EWR  BQN     208     10.864     -43   297
#> 10  EWR  BTV     306     12.186     -41   931
#> # ... with 214 more rows

#- derived columns: partition air_time into 5 categories (check the NA row too)
by_air_time = group_by(flights, air_time2 = cut(air_time, 5)) # added column
summarize(by_air_time,
  max.delay = max(arr_delay, na.rm=TRUE),
  avg.delay = mean(arr_delay, na.rm=TRUE),
  min.delay = min(arr_delay, na.rm=TRUE),
  count = n() )
#> # A tibble: 6 x 5
#>   air_time2 max.delay avg.delay min.delay count
#>   <fctr>    <dbl>    <dbl>    <dbl> <int>
#> 1 (19.3,155]    1127      7.63      -68 215127
#> 2 (155,290]     915      7.72      -75 64954
#> 3 (290,425]    1007      2.45      -86 46558
#> 4 (425,560]      92     48.00       -4     6
#> 5 (560,696]    1272     -1.37     -70    701
#> 6      NA      NA      NaN      NA   9430

```

### 2.2.1 Your Turn

#### Your Turn #1 : group\_by

Which plane (tailnum) has the worst on-time record?

## 2.3 Counting

We often need to count the number of observations in each group. It is so frequently needed, that `dplyr` has included some shortcuts with `n()`, `tally()` and `count()`

1. use `summarize()` with `n()` (must use *grouped* data)

```
#- tally(.) is same as summarize(., n=n())
summarize(by_origin, n=n())
#> # A tibble: 3 × 2
#>   origin      n
#>   <chr>   <int>
#> 1   EWR 120835
#> 2   JFK 111279
#> 3   LGA 104662
```

## 2. use tally() (must use grouped data)

```
tally(by_origin, sort=TRUE)
#> # A tibble: 3 × 2
#>   origin      n
#>   <chr>   <int>
#> 1   EWR 120835
#> 2   JFK 111279
#> 3   LGA 104662
```

## 3. use count() (don't use grouped data)

```
#- count(., colname) is same as group_by(., colname) %>% tally()
count(flights, origin)
#> # A tibble: 3 × 2
#>   origin      n
#>   <chr>   <int>
#> 1   EWR 120835
#> 2   JFK 111279
#> 3   LGA 104662
```

### 2.3.1 Counts over multiple variables

```
#- Count for each route (origin and destination)
count(flights, origin, dest)
#> Source: local data frame [224 × 3]
#> Groups: origin [?]
#>
#>   origin dest      n
#>   <chr> <chr> <int>
#> 1   EWR  ALB    439
#> 2   EWR  ANC      8
#> 3   EWR  ATL   5022
#> 4   EWR  AUS    968
#> 5   EWR  AVL    265
#> 6   EWR  BDL    443
#> 7   EWR  BNA   2336
#> 8   EWR  BOS   5327
#> 9   EWR  BQN    297
#> 10  EWR  BTV    931
#> # ... with 214 more rows

#- Count for each route by month
count(flights, origin, dest, month)
```

```
#> Source: local data frame [2,313 x 4]
#> Groups: origin, dest [?]
#>
#>   origin dest month     n
#>   <chr> <chr> <int> <int>
#> 1    EWR  ALB     1     64
#> 2    EWR  ALB     2     58
#> 3    EWR  ALB     3     57
#> 4    EWR  ALB     4     13
#> 5    EWR  ALB     5     59
#> 6    EWR  ALB     6     34
#> 7    EWR  ALB     7     15
#> 8    EWR  ALB     8     20
#> 9    EWR  ALB     9     20
#> 10   EWR  ALB    10      1
#> # ... with 2,303 more rows
```

### 2.3.2 Plotting the counts

Get the monthly counts

```
(monthly = count(flights, origin, month))
#> Source: local data frame [36 x 3]
#> Groups: origin [?]
#>
#>   origin month     n
#>   <chr> <int> <int>
#> 1    EWR     1  9893
#> 2    EWR     2  9107
#> 3    EWR     3 10420
#> 4    EWR     4 10531
#> 5    EWR     5 10592
#> 6    EWR     6 10175
#> 7    EWR     7 10475
#> 8    EWR     8 10359
#> 9    EWR     9  9550
#> 10   EWR    10 10104
#> # ... with 26 more rows
```

Notice that `count()` creates the column named `n` (integer).

```
#- (left) Bar Plot
ggplot(monthly) + geom_col(aes(x=month, y=n, fill=origin))
# ggplot(flights) + geom_bar(aes(x=month, fill=origin)) # alternative

#- (right) Line Plot
ggplot(monthly) + geom_line(aes(x=month, y=n, col=origin))
```





### 2.3.3 Your Turn

#### Your Turn #2 : Thinking about plots

1. Are the plots better than the table?
2. Which plot do you think is better?
3. How can the plots be improved?

### 2.3.4 Additional arguments

Check out the help for the `count()` function: `?count` There are two additional arguments:

- `wt` gives a weighted sum (instead of plain count)
- `sort` will sort from largest to smallest

```
#-- total arrival delay by flight number
count(flights, flight, wt=arr_delay, sort=TRUE )
#> # A tibble: 3,844 × 2
#>   flight      n
#>   <int> <dbl>
#> 1   4131 12989
#> 2    527 11694
#> 3   4333 11433
#> 4    415 11390
#> 5   4224 10204
#> 6   4543 10148
#> 7   1161 10132
#> 8    985 10014
#> 9   2042  9777
#> 10  4204  9478
#> # ... with 3,834 more rows
```

### 2.3.5 Other types of count

Some useful counts are

- count the number of distinct items with `n_distinct()`

```

summarize(by_origin, n_flights=n(), n_dests=n_distinct(dest))
#> # A tibble: 3 × 3
#>   origin n_flights n_dests
#>   <chr>   <int>   <int>
#> 1 EWR     120835     86
#> 2 JFK     111279     70
#> 3 LGA     104662     68

```

- count the number of non-missing (not-NA) values

```

summarize(by_origin, n_flights=n(), n_missing = sum(is.na(dep_time)),
          n_not_missing=sum(!is.na(dep_time)))
#> # A tibble: 3 × 4
#>   origin n_flights n_missing n_not_missing
#>   <chr>   <int>   <int>   <int>
#> 1 EWR     120835     3239     117596
#> 2 JFK     111279     1863     109416
#> 3 LGA     104662     3153     101509

```

## 2.3.6 Your Turn

### Your Turn #3 : counting

1. How many flights does the plane with the worst on-time record have?
2. Which plane (tailnum) has made the most flights?
3. Which plane (tailnum) has flown the most distance?

## 2.4 Chaining

Multiple operations can be chained together with the `%>%` operator (pronounced as *then*). Technically, it performs `x %>% f(y) -> f(x, y)`. This lets you focus on the verbs, or actions you are performing.

```

#- group then summarize then filter then arrange
by_dest = group_by(flights, dest)
delay = summarize(by_dest,
                  count = n(),
                  avg_delay = mean(arr_delay, na.rm=TRUE))
delay2 = filter(delay, count > 20)
arrange(delay2, desc(avg_delay))
#> # A tibble: 97 × 3
#>   dest count avg_delay
#>   <chr> <int>   <dbl>
#> 1 CAE    116    41.8
#> 2 TUL    315    33.7
#> 3 OKC    346    30.6
#> 4 JAC     25    28.1
#> 5 TYS    631    24.1
#> 6 MSN    572    20.2
#> 7 RIC   2454    20.1

```

```

#> 8    CAK    864    19.7
#> 9    DSM    569    19.0
#> 10   GRR    765    18.2
#> # ... with 87 more rows

flights %>%
  group_by(dest) %>%
  summarize( count = n(),
             avg.delay = mean(arr_delay, na.rm=TRUE)) %>%
  filter(count > 20) %>%
  arrange(desc(avg.delay))
#> # A tibble: 97 × 3
#>   dest count avg.delay
#>   <chr> <int>   <dbl>
#> 1  CAE    116    41.8
#> 2  TUL    315    33.7
#> 3  OKC    346    30.6
#> 4  JAC     25    28.1
#> 5  TYS    631    24.1
#> 6  MSN    572    20.2
#> 7  RIC   2454    20.1
#> 8  CAK    864    19.7
#> 9  DSM    569    19.0
#> 10 GRR    765    18.2
#> # ... with 87 more rows

```

## Your Turn #4 : Chaining

Use chaining to redo:

1. How many flights does the plane with the worst on-time record have?
2. Which plane (tailnum) has made the most flights?
3. Which plane (tailnum) has flown the most distance?

## 2.5 Multiple grouping levels

Notice how each operation strips away a grouping level.

```

#- grouped by: year, month, day
(daily <- group_by(flights, year, month, day))
#> Source: local data frame [336,776 x 19]
#> Groups: year, month, day [365]
#>
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>       <dbl>   <int>
#> 1  2013     1     1     517           515         2       830
#> 2  2013     1     1     533           529         4       850
#> 3  2013     1     1     542           540         2       923
#> 4  2013     1     1     544           545        -1      1004
#> 5  2013     1     1     554           600        -6       812
#> 6  2013     1     1     554           558        -4       740

```

```

#> 7  2013      1      1      555      600      -5      913
#> 8  2013      1      1      557      600      -3      709
#> 9  2013      1      1      557      600      -3      838
#> 10 2013      1      1      558      600      -2      753
#> # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>

#- grouped by: year, month
(per_day <- summarise(daily, flights = n()))
#> Source: local data frame [365 x 4]
#> Groups: year, month [?]
#>
#>   year month   day flights
#>   <int> <int> <int>   <int>
#> 1  2013     1     1     842
#> 2  2013     1     2     943
#> 3  2013     1     3     914
#> 4  2013     1     4     915
#> 5  2013     1     5     720
#> 6  2013     1     6     832
#> 7  2013     1     7     933
#> 8  2013     1     8     899
#> 9  2013     1     9     902
#> 10 2013     1    10     932
#> # ... with 355 more rows

#- grouped by: year
(per_month <- summarise(per_day, flights = sum(flights)))
#> Source: local data frame [12 x 3]
#> Groups: year [?]
#>
#>   year month flights
#>   <int> <int>   <int>
#> 1  2013     1  27004
#> 2  2013     2  24951
#> 3  2013     3  28834
#> 4  2013     4  28330
#> 5  2013     5  28796
#> 6  2013     6  28243
#> 7  2013     7  29425
#> 8  2013     8  29327
#> 9  2013     9  27574
#> 10 2013    10  28889
#> 11 2013    11  27268
#> 12 2013    12  28135

#- grouped by: nothing (i.e., this is not grouped data)
(per_year <- summarise(per_month, flights = sum(flights)))
#> # A tibble: 1 x 2
#>   year flights
#>   <int>   <int>

```

```
#> 1 2013 336776
```

If you want to remove the grouping, use `ungroup()` function.

### 2.5.1 Your Turn

#### Your Turn #5 : Multiple groups

1. Find the top 5 routes (`origin`, `dest`), in terms of number of flights.
2. Which route (`origin`, `dest`) is most often delayed by more than 10 minutes? Are infrequent routes a concern? If so, what could we do about it?
3. Find the top 3 destinations (`dest`) for each origin (`origin`).

## 2.6 Grouped Mutate and Filter

The last exercise (find the top 3 destinations for each origin) requires a *grouped filter*. That is, perform filtering *within* each group separately.

A *grouped mutate* can calculate standardizations per group

```
#- proportion of carrier at each dest
flights %>%
  count(dest, carrier) %>%           # still grouped by dest
  mutate(total=sum(n), p=n/sum(n)) %>% # grouped mutate sum(n) is by group
  arrange(desc(total))               # arrange by most freq dest
#> Source: local data frame [314 x 5]
#> Groups: dest [105]
#>
#>   dest carrier      n total      p
#>   <chr>   <chr> <int> <int>   <dbl>
#> 1  ORD     9E   1056 17283 6.11e-02
#> 2  ORD     AA   6059 17283 3.51e-01
#> 3  ORD     B6    905 17283 5.24e-02
#> 4  ORD     EV     2 17283 1.16e-04
#> 5  ORD     MQ   2276 17283 1.32e-01
#> 6  ORD     OO     1 17283 5.79e-05
#> 7  ORD     UA   6984 17283 4.04e-01
#> 8  ATL     9E     59 17215 3.43e-03
#> 9  ATL     DL  10571 17215 6.14e-01
#> 10 ATL     EV   1764 17215 1.02e-01
#> # ... with 304 more rows
```

To do a *grouped filter* or *mutate*, the data frame must be grouped with `group_by()` or `count()`. Functions that work most naturally in grouped mutates and filters are known as *window functions*.

## 2.7 Window Functions

A *window function* is a variation on an aggregation function. Where an aggregation function, like `sum()` and `mean()`, takes `n` inputs and return a single value, a window function returns `n` values. The output of a window function depends on all its input values, so window functions don't include functions that work element-wise, like `+` or `round()`. Window functions include variations on aggregate functions, like `cumsum()` and `cummean()`, functions for ranking and ordering, like `rank()`, and functions for taking offsets, like `lead()` and `lag()`. The Z-score can be considered a window function.

More description of some window functions and their use can be found:

- [data transform cheatsheet](#)

- [Window function vignette](#)

The rank family (`?min_rank`) of function can help you not get burned by the ordering. For example, we wanted the top 3 destinations (`dest`) for each origin (`origin`).

```
#- Note: still grouped by origin. So slice operates on group. Must be sorted!
flights %>% count(origin, dest, sort=TRUE) %>% slice(1:3)
#> Source: local data frame [9 x 3]
#> Groups: origin [3]
#>
#>   origin dest      n
#>   <chr> <chr> <int>
#> 1   EWR   ORD   6100
#> 2   EWR   BOS   5327
#> 3   EWR   SFO   5127
#> 4   JFK   LAX  11262
#> 5   JFK   SFO   8204
#> 6   JFK   BOS   5898
#> 7   LGA   ATL  10263
#> 8   LGA   ORD   8857
#> 9   LGA   CLT   6168

#- incorrect with sort=TRUE
flights %>% count(origin, dest) %>% slice(1:3)
#> Source: local data frame [9 x 3]
#> Groups: origin [3]
#>
#>   origin dest      n
#>   <chr> <chr> <int>
#> 1   EWR   ALB    439
#> 2   EWR   ANC      8
#> 3   EWR   ATL   5022
#> 4   JFK   ABQ    254
#> 5   JFK   ACK    265
#> 6   JFK   ATL   1930
#> 7   LGA   ATL  10263
#> 8   LGA   AVL     10
#> 9   LGA   BGR    375

#- always correct with rank:
flights %>% count(origin, dest) %>% filter(min_rank(-n) <= 3) %>%
```

```
arrange(origin, desc(n))
#> Source: local data frame [9 x 3]
#> Groups: origin [3]
#>
#>   origin dest      n
#>   <chr> <chr> <int>
#> 1   EWR   ORD   6100
#> 2   EWR   BOS   5327
#> 3   EWR   SFO   5127
#> 4   JFK   LAX  11262
#> 5   JFK   SFO   8204
#> 6   JFK   BOS   5898
#> 7   LGA   ATL  10263
#> 8   LGA   ORD   8857
#> 9   LGA   CLT   6168
```